

AWS re:Invent

DECEMBER 1 - 5, 2025 | LAS VEGAS, NV

CNS 359

Build, deploy, and operate agentic architectures on AWS Serverless

Heeki Park

(he/him)

Principal Solutions Architect

AWS

Dhiraj Mahapatro

(he/him)

Principal Solutions Architect

AWS

Agenda

- Designing single-agent systems
- Applying serverless to agentic patterns
- Developing agents
- Scaling with open standards
- Deploying agents
- Transitioning to multi-agent systems
- Baking in security and governance for deployments



Assumptions



Level 300
session



Principles
first

Designing single-agent systems

Starting with a customer issue

1

2

3


I was charged twice for my order, only received one of the two items in the order, and I need it tomorrow for a gift.

Can you assist in rectifying and expediting?



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Starting with serverless APIs



Considering customer experience




Challenges:


1. Responding to customer requests requires **multiple lookups in disparate systems**.
2. Customers need **faster time to resolution**.
3. Support agents need **faster guidance** when responding to customer calls or chat requests.

Can we do better?


Adding LLMs to applications



Providing more context through search



Providing more context **dynamically**




What does an enhanced prompt look like?

ADDITIONAL CONTEXT


```
{  
  "system_prompt": "You are a helpful assistant..",  
  "tools": [],  
  "context": "...",  
  "conversation_history": "...",  
  "user_info": "...",  
  "human_prompt": "..."  
}
```




Providing context dynamically **through APIs**




Transitioning to agents




Why agents?



Decomposing the agentic monolith



Refining domain agents




Initial tool functionality
based on existing API
capabilities


Extended tool functionality
based on model inference

Applying serverless to agentic patterns


Serverless with Tool use




Serverless with Tool use




Serverless with Tool use




Serverless with Tool use



Tool use in AWS Step Functions



Emerging pattern



Tool use in Lambda or Step Functions



Converse/Invoke API
Amazon Bedrock




Imperative code to
act on behalf of
LLM



Predefined pathways
& workflows

Agentic AI with Strands Agents SDK



Developing agents

< Hello, World! >

Strands Agents

Strands Agents is an open-source SDK for building agents using just a few lines of code



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Simple Strands agent

```
from strands import Agent
from strands_tools import calculator

agent = Agent(tools=[calculator])
response = agent("What is 80 / 4?")
```



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Order support using Strands Agents SDK

```
1 import json
2 import boto3
3 from datetime import datetime
4 from strands import Agent, tool
5 from strands.models import BedrockModel
6
```

Import from SDK

```
7 @tool
8 > def get_order_details(order_id: str) -> dict: ...
23
24 @tool
25 > def get_invoice(order_id: str) -> dict: ...
62
63 @tool
64 > def start_refund_process(order_id: str, reason: str) -> dict: ...
97
```

Define methods as tools

```
98 session = boto3.Session(
99     region_name='us-east-1'
100 )
101
102 bedrock_model = BedrockModel(
103     model_id="us.anthropic.claude-4-5-haiku-20251001-v1:0",
104     max_tokens=5000,
105     boto_session=session,
106 )
107
```

Use Bedrock as model provider

```
108 # Initialize the Strands agent with tools
109 agent = Agent(
110     model=bedrock_model,
111     tools=[get_order_details, get_invoice, start_refund_process]
112 )
113
```

Create agent with model and tools


```
114 agent["I was charged twice for my recent order"]
```

Invoke agent with user prompt



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Benefits of using an agentic framework



In-built error handling,
retries, other best practices




Model and compute
agnostic




Supports dynamic pathways
(agentic loop)
& workflows

Building an agent is the start, production scale is the goal




Scaling agents and tools with
open standards

Agents invoking remote tools



Agent-tool integrations without a standard




Agent builders are responsible to use different interfaces to integrate with tools


Restricted scope for tools vendors or cater to multiple interfaces.



Agent integrations with Model Context Protocol (MCP)



Standardizing tools integration with MCP



Using MCP client in a Strands agent

```
web_search_agent.py > ...
1  from mcp.client.streamable_http import streamablehttp_client
2  from strands import Agent
3  from strands.tools.mcp.mcp_client import MCPClient
4
5  streamable_http_mcp_client = MCPClient(lambda: streamablehttp_client("https://server.smithery.ai/exa/mcp?api_key=<api-key>&profile=<profile>"))
6
7  # Create an agent with MCP tools
8  with streamable_http_mcp_client:
9      # Get the tools from the MCP server
10     tools = streamable_http_mcp_client.list_tools_sync()
11
12     # Create an agent with these tools
13     agent = Agent(
14         system_prompt="You are an assistant that searches the web on asked topics and provides a concise result",
15         tools=tools
16     )
17
18     response = agent(["What are the top 5 best rated cities in USA grouped by breweries? Provide answers in tabular form."])
```

Import Strands agent and MCP Client

Initialize MCP client

List remote tools and provide to agent

```
(.venv) ➜ my_agent python -u web_search_agent.py
I'll search for information about the top-rated cities in the USA based on their breweries
Tool #1: web_search_exa
Based on my research, I can provide you with a comprehensive table of the top 5 best-rated
USA Today's 10Best rankings, and Beer & Brewing magazine's reader's choice awards.

## Top 5 Best Rated Cities in USA for Breweries

| Rank | City | State | Key Strengths | Notable Features |
|-----|-----|-----|-----|-----|
| **1** | **Denver** | Colorado | Over 100 breweries, brewpubs, and taprooms; Home to Great American Beer Festival |
| **2** | **Portland** | Oregon | "Beervana" - Over 70 breweries; Innovation hub for craft beer |
```

Building an MCP server

```
	mc_server_demo.py > ...
1  from fastmcp import FastMCP
2
3  mcp = FastMCP(name="Basic Maths MCP Server")
4
5  @mcp.tool
6  def add(a: int, b: int) -> int:
7      """Add two numbers"""
8      return a + b
9
10 if __name__ == "__main__":
11     mcp.run(transport="http", host="127.0.0.1", port=8000, stateless_http=True)
```

Import and create FastMCP

Create a tool which will be used by an MCP Client remotely

Run the server

```
(.venv) > my_agent fastmcp dev mcp_server_demo.py
Starting MCP inspector...
Proxy server listening on localhost:6277
Session token: 8b967c879477a5b25ee25456f6bfad0271149c875fedfed0f50d9568ced0a997
Use this token to authenticate requests or set DANGEROUSLY OMIT_AUTH=true to disable auth


MCP Inspector is up and running at:
http://localhost:6274/?MCP_PROXY_TOKEN=8b967c879477a5b25ee25456f6bfad0271149c875fedfed0f50d9568ced0a997

Opening browser...
New STDIO connection request
```




© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Agents invoking remote **agents**




Inter-agent without a standard



Challenges:

- Scope creep
- Run tools from a different domain
- Ownership leakage
- Continuous revision


Inter-agent with MCP



- [Improvement] Better than no standards
- [Challenges]
 - Payments agent is used as a tool
 - Order agent must discover about Payments agent's tools
 - Payments agent must act as an MCP Server & MCP Client

Introducing Agent-to-Agent (A2A)


A standardized protocol for AI agents to discover capabilities, exchange tasks, and collaborate on complex workflows



Agent Card

- name/version/desc.
 - skills
 - capabilities
 - supported modalities
 - auth requirements

Inter-agent with A2A




Improvements:

- Encourages to build specialized agents
- Payments agent doesn't leak domain knowledge
- Order agent discovers skills and capabilities of Payments agent from its agent card
- A2A supports sync and async communications



Event-driven with A2A and MCP



Benefits of open standards and EDA



Loose coupling
with intelligent
coordination



Context-aware
event handling



Scalable agent
orchestration



Asynchronous
intelligence



Resilient
workflows that
are observable
and auditable



Dynamic agent
subscription

Deploying agents and tools
on Serverless compute

Deploying with AWS Serverless compute



AWS Lambda




Amazon Elastic
Container Service (ECS)



Amazon Bedrock
AgentCore


Agentic patterns for serverless compute




Deploying agents to ECS Fargate




Deploying agents to Lambda




Deploying tools (MCP Servers)




A2A and MCP on Lambda functions




Considering serverless agents




Deploying agents to AgentCore Runtime




Deploying MCP servers to AgentCore Runtime



AgentCore Gateway with Lambda as tool



A2A and MCP on AgentCore



Transitioning to multi-agent systems

Decomposed agents

Orders Agent



Inventory Agent



Payments Agent



Support Agent



Decomposed agents



Orders Agent



Inventory Agent




Payments Agent




Support Agent


Orchestrating agent work



Scaling multi-agent systems



Starting with single-agent architecture



Pros:

- simpler to implement and debug
- easier to maintain and modify
- clearer reasoning traces
- *possibly* less computational overhead

Cons:

- limited by single context window
- can get stuck in reasoning loops
- could lack specialization, if required
- may struggle with complex, multi-step workflows

Building multi-agent architecture




Pros:

- natural workflow modeling
- clear handoffs and error localization
- easy to optimize individual stages

Cons:

- limited to serial behavior, no benefits of parallelization
- serial dependencies
- errors amplified through the chain
- inflexible for non-linear processes

Building multi-agent architecture




Pros:

- task decomposition, which allows for specialization and scalability
- composability allows for new agents
- maps more naturally to business process
- fault isolation

Cons:

- communication overhead
- orchestrator could become bottleneck
- error handling adds complexity
- could be rigid for dynamic problem-solving

Building multi-agent architecture



Pros:


- highly adaptive, as agents self-organize based on task assignment
- has emergent problem-solving capabilities
- fault tolerant
- natural load balancing

Cons:


- difficult to predict behavior
- hard to debug and trace decisions
- complex coordination across agents
- requires agent discovery mechanisms

Baking in security and governance for deployments

Considering security



Considering governance



Takeaways

The **same** design and architecture principles **continue to apply** with agents

Agents can be a way to **bridge data and glean insights**

Serverless can be a way to **accelerate the process** for building your agents

Serverless can help you **scale your agents** using standards and managed services



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Key considerations

Be clear about your use cases. Fully understand where generative AI is appropriate.

Clearly document measurement metrics. Know what success looks like before you start.

Secure and instrument from the beginning. Security and observability cannot be after-thoughts.

Build and be scrappy. Lots to be uncovered during prototyping.

Document and provide feedback. Platform engineering needs to understand what they need to build and service for your developers.



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Resources



<https://s12d.com/riv25-cns359>



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Other sessions

[DEV415] Building Scalable, Self-Orchestrating AI Workflows with A2A and MCP

Monday, December 1 @ 1:30pm | MGM, Level 3

[CNS425] Building production-ready Agentic AI architectures with AWS serverless

Monday, December 1 @ 5:30pm | MGM | Level 3 | Chairman's 356

Tuesday, December 2 @ 1:00pm | Mandalay Bay | Level 2 South | Reef C

[CNS360] Implementing security best practices for serverless applications

Tuesday, December 2 @ 3pm | Mandalay Bay | Level 2 South, Oceanside C

[CNS361] Scaling Serverless with platform engineering: A blueprint for success

Wednesday, December 3 @ 4:30pm | Mandalay Bay, Level 2 South, Oceanside C

[CNS428] Implementing Human-in-the-loop Controls for Multi-Agent AI Systems

Thursday, December 4 @ 2:30pm | Mandalay Bay, Lightning Theater



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Additional resources



<https://s12d.com/RIV2025>

Continue your AWS serverless learning
PowerTools for AWS Lambda
Serverless Land patterns

 **Serverless Land** [Services](#) ▾ [Learn](#) ▾ [Example code](#) ▾ [EDA](#) ▾ [Search](#)

Welcome to **Serverless Land**

This site brings together the latest information, blogs, videos, code, and learning resources for AWS Serverless. Learn to use and build apps that scale automatically on low-cost, fully-managed serverless architecture.



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.





Thank you!

Heeki Park

 linktr.ee/heekipark

Dhiraj Mahapatro

 [@dmahapatro](https://www.linkedin.com/in/dmahapatro)

 linktr.ee/dmahapatro

Please complete the session
survey in the mobile app